

# GO in the Distributed Systems world

Alberto Giunta

Department of Computer Science  
University of Bologna

October 30, 2017

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

# Let's GO!

## The What, Where and Why of GO

### What

GO is a compiled, concurrent, garbage-collected, statically typed, open source programming language, which aims to be the *next* C.

GO is about language design in the service of software engineering.

### Where and Why

Created at Google in 2007 in order to solve the problems they were facing with their *really* large distributed and computation-based architecture

### Put in other words

"All three of the language designers cited their shared dislike of C++'s complexity as a primary motivation for designing a new language" [4]

# Let's GO!

## The What, Where and Why of GO

- It was built and designed with a few requirements in mind:
  - It must work at scale (as Java and C++)
  - It must be familiar, productive and readable
  - It must support networking and multiprocessing
- When implemented, it was decided to follow this path:
  - Start from C
  - Remove its complex parts
  - add first class support to solutions for modern problems (i.e. concurrency)

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

# Let's GO!

## PROs and CONs of GO

### PROs:

- C family (imperative, braces)
- Statically typed
- Quick compilation times
- First-class support to concurrency
- Well suited language for networking projects, distributed functions or services
- Everything is a function

# Let's GO!

## PROs and CONs of GO

### CONs:

- Lacks of generics
- No constructors
- No exceptions
- No annotations
- No classes (they can be achieved through the use of Structs though)
- No inheritance (GO strictly follows the "Composition over Inheritance" law)



# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - **Install GO on your machine**
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

# It's time to GO!

Let's install GO on your machine

- Open IntelliJ and install the GO plugin
- [Click here](#) to download the GO binaries and install them on your local machine
- Set the environment variables in your \*nix console:
  - \$ export GOPATH=\$HOMEgo
  - \$ export PATH=\$PATH:\$GOPATHbin
- Let's run our first GO program hello.go!

```
1 package main
2
3 import "fmt" // the standard library
4
5 func main() {
6     fmt.Printf("Hello, world.\n")
7 }
```

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - **Basic syntax and constructs**
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

# Structs

## Characteristics of Structs

- A Struct is a typed collection of named fields
- They're useful for grouping data together to form records
- Structs and Interfaces, together, make it possible to have OOP in GO (i.e. you can see them as Java classes, with field and methods)

# Structs

```
1 package main
2 import "fmt"
3
4 type person struct {
5     name string
6     age int
7 }
8
9 func main() {
10     fmt.Println(person{"Bob", 20}) // create a new struct.
11     // the & creates a pointer to the struct
12     fmt.Println(&person{name: "Ann", age: 40})
13     s := person{name: "Sean", age: 50}
14     fmt.Println(s.name) // access struct fields with a dot.
15     sp := &s
16     fmt.Println(sp.age) // struct pointers are deferentiated
17     sp.age = 51 // Structs are mutable.
18     fmt.Println(sp.age)
19 }
```

# Pointers

## Characteristics of Pointers

- They basically behave just like C pointers do: `*` and `&` respectively represent a pointer to a value in memory, and its memory address
- Pointers are available for all types, and the pointer-to-T type is denoted `*T`
- There's no pointer arithmetic though

```
1 p := Vertex{1, 2} // p is a Vertex
2 q := &p           // q is a pointer to a Vertex
3 r := &Vertex{1, 2} // r is also a pointer to a Vertex
4
5 // The type of a pointer to a Vertex is *Vertex
6
7 var s *Vertex = new(Vertex) // new creates a pointer to a new struct instance
```

# Interfaces

## Characteristics of Interfaces

Similar to Structs, but instead of declaring fields, they contain a "method set", which is a list of methods that a type must have in order to "implement" the interface (hence they behave just like interfaces behave in languages like java)

```
1 // interface declaration
2 type Awesomizer interface {
3     Awesomize() string
4 }
5
6 // types do *not* declare to implement interfaces
7 type Foo struct {}
8
9 // instead, types implicitly satisfy an interface if they implement all required methods
10 func (foo Foo) Awesomize() string {
11     return "Awesome!"
12 }
```

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system



# CSP

## Communicating Sequential Processes

### What is CSP

- GO is built upon CPS (Communicating Sequential Processes)
- First described in a 1977 paper by Tony Hoare
- Formal language for describing patterns of interaction in concurrent systems
- It is a member of the family of mathematical theories of concurrency known as process algebras, based on message passing via channels
- Independent executing processes that communicate by passing messages

# Differences between CSP and the Actor Model

They're both about concurrent processes that exchange messages, but with some key differences in their primitives:

## CSP

- CSP processes are anonymous
- CSP message-passing fundamentally involves a rendezvous between the processes involved in sending and receiving the message, i.e. the sender awaits for a response message sent by the receiver after it received the first message
- CSP uses explicit channels for message passing

# Differences between CSP and the Actor Model

## Actor Model

- Actors have identities
- Message-passing in actor systems is fundamentally asynchronous, i.e. message transmission and reception do not have to happen at same time, and senders may transmit messages before receivers are ready to accept them.
- Actor systems transmit messages to named destination actors

# Goroutine

## What is a Goroutine

- A Goroutine is a function that is capable of running concurrently with other functions
- In order to create a Goroutine we use the keyword `go` followed by a function invocation
- Goroutines are lightweight and we can easily create thousands of them

## Use case of a Goroutine

- Normally when we invoke a function our program will execute all the statements in a function and then return to the next line following the invocation.
- With a goroutine we return immediately to the next line and don't wait for the function to complete

# Goroutine

## Using Goroutines

```
1 package main
2 import "fmt"
3
4 func printAsync(from string) {
5     for i := 0; i < 3; i++ {
6         fmt.Println(from, ":", i)
7     }
8 }
9
10 func main() {
11     printAsync("direct")           // function executed synchronously
12
13     go printAsync("goroutine")     // function executed asynchronously with goroutine
14
15     go printAsync(msg string) {
16         fmt.Println(msg)          // anonymous function inside goroutine
17     }("going")
18
19     var input string
20     fmt.Scanln(&input)
21     fmt.Println("done")
22 }
```

# Channel

## The good ol' problem of sharing memory in a concurrent program

- GO enforces the concept of "Do not communicate by sharing memory; instead, share memory by communicating"
- Go encourages a different approach in which shared values are passed around on channels and, in fact, never actively shared by separate threads of execution
- Only one Goroutine has access to the value at any given time. Data races cannot occur, by design.
- Channels are the pipes that connect concurrent Goroutines. You can send values into channels from one Goroutine and receive those values into another Goroutine.

# Channel

## Channels

- A channel provides a mechanism for concurrently executing functions to communicate by sending and receiving values of a specified element type
- Channels provide a way for two Goroutines to communicate with one another and synchronize their execution
- A channel type is represented with the keyword `chan` followed by the type of the things that are passed on the channel
- By default sends and receives block until both the sender and receiver are ready.

# Channels

A quick sample of how channels work

## Use case of a Channel

The `<-` (left arrow) operator is used to receive messages on a certain channel or to receive a value from a channel and put it in a variable

```
1 ch := make(chan int) // create a channel of type int
2 ch <- 42              // Send a value to the channel ch.
3 v := <-ch             // Receive a value from ch
4 close(ch)            // closes the channel (only sender should close)
```



# Channel

## Buffered Channel

It's also possible to pass a second parameter to the `make` function when creating a channel (i.e. create a channel of capacity of 1):

```
c := make(chan int, 1)
```

## Difference between a Buffered and Non-buffered Channel

- Normally channels are synchronous; both sides of the channel will wait until the other side is ready.
- A buffered channel is asynchronous; sending or receiving a message will not wait unless the channel is already full.

# Select

## Select

- The Select statement lets you wait on multiple channel operations.
- Combining Goroutines and Channels with Select is a powerful feature of Go
- Each channel will receive a value after some amount of time, to simulate e.g. blocking RPC operations executing in concurrent Goroutines

```
1 // select blocks on multiple channel operations
2 // if one unblocks the corresponding case is executed
3 func doStuff(channelOut, channelIn chan int) {
4     select {
5         case channelOut <- 42:
6             fmt.Println("Write 42 into channelOut")
7         case x := <- channelIn:
8             fmt.Println("Save from channelIn into x")
9         case <-time.After(time.Second * 1):
10             fmt.Println("Timeout received. Not going to accept any other message here.")
11     }
12 }
```

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

# Es 1: Ping Pong

We will now try to build a simple program that simulates the Ping Pong game, by using Goroutines and Channels.

```
1 package main
2 import "fmt"
3
4 func player(playerName string, ch chan string, rounds int) {
5     // TODO create the player behaviour, by implementing a loop where
6     // the player initially waits for a message,
7     // then prints its own name
8     // and then sends another message on the channel to unlock who's waiting for a message
9 }
10
11 func main() {
12     fmt.Println("Welcome to the Ping Pong Game!")
13     rounds := 10
14
15     // TODO create a new channel that will be used to transmit strings
16
17     // TODO create two goroutines where you create two players, pinger and ponger
18
19     // TODO send an initial message "init" on the channel to start the players
20
21     var input string
22     fmt.Scanln(&input)
23 }
```

# Ex 1: Ping Pong

## Solution

```
1 package main
2 import "fmt"
3
4 func player(playerName string, ch chan string, rounds int) {
5     for i := 0; i < rounds; i++ {
6         <-ch
7         fmt.Println(playerName)
8         ch <- "unlock"
9     }
10 }
11
12 func main() {
13     fmt.Println("Welcome to the Ping Pong Game!")
14
15     rounds := 10
16
17     var sharedChannel chan string = make(chan string)
18
19     go player("ping", sharedChannel, rounds)
20     go player("pong", sharedChannel, rounds)
21
22     sharedChannel <- "init"
23
24     var input string
25     fmt.Scanln(&input)
26 }
```

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

# searchUtils.go

## What are we going to build

- We're going to build a distributed parallel system that will retrieve some information from different sources in different locations, just like Google does
- For example, we query "foo" and the system will give us results regarding Websites, Images and Videos about our query "foo"
- Clearly, we're not Google ourselves, so we're going to mock the data, and just implement the basic logic behind it all

## How is the code organized

- The code is organized in two different packages:
  - `google_launchers` contains all the main programs, and each launches the system using a different function and logic
  - `google` contains all the files regarding the system's logic. This is where the magic happens and where we're going to write our algorithms

# Parallel Computation and Systems

## Parallel computation

A computation is parallel whenever the temporal context is the same for all computational process

## Parallel system

A parallel system is a computational system performing parallel computations



# Distributed Computation and Systems

## Distributed computation

A computation is distributed whenever at least two computational processes have a different spatial context

## Distributed system

A distributed system is a computational system performing distributed computations




Figure: Distributed parallel computing:  $S \neq S'$ , same  $T$

# Quality of Service

Since we're going to iterate quite a lot of times over the next exercise, it's important to understand why.

A big part of the problems we face while building a distributed system come from the importance of QoS

## QoS

*Quality of service (QoS) is the description or measurement of the overall performance of a service, such as a telephony or computer network or a cloud computing service, particularly the performance seen by the users of the network.[1]*

It the next few slides we're going to implement various versions of our systems, and we'll try to raise the QoS level at each iteration[2]

# searchUtils.go

## Already given and functioning code

SearchUtils.go is an helper class that contains the code regarding the data in the different repositories and some other simple data types

```
1 // function that will be passed around to be executed in all the different ways (serial, parallel etc.)
2 type SearchFunc func(query string) Result
3
4 // the resulting data type of a Search
5 type Result struct {
6     Title, URL string
7 }
8
9 // the data type we'll use to display the results on screen
10 type Response struct {
11     Results []Result
12     Elapsed time.Duration
13 }
14
15 // function that will be passed around to be executed in all the different ways (serial, parallel, with timeout, distributed)
16 func FakeSearch(kind, title, url string) SearchFunc {
17     return func(query string) Result {
18         time.Sleep(time.Duration(rand.Intn(100)) * time.Millisecond) // artificial delay for instructional reasons
19         return Result{
20             Title: fmt.Sprintf("%s(%q): %s", kind, query, title),
21             URL:    url,
22         }
23     }
24 }
```

# searchUtils.go

Already given and functioning code

The functions that will be used in the various types of data fetching processes

```
1 // serial, parallel, with timeout
2 var (
3     Web = FakeSearch("web", "Distributed Systems are the best!", "http://apice.unibo.it/xwiki...")
4     Image = FakeSearch("image", "But what's really a distributed system?", "https://www.unibo.it/uniboweb...")
5     Video = FakeSearch("video", "Cool video about distributed systems", "https://www.youtube.com/watch?v=...")
6 )
7
8 // distributed
9 var (
10     replicatedWebFakeSearch = FirstFetchAmong(Web1, Web2)
11     replicatedImageFakeSearch = FirstFetchAmong(Image1, Image2)
12     replicatedVideoFakeSearch = FirstFetchAmong(Video1, Video2)
13
14     Web1 = FakeSearch("web1", "Distributed Systems are the best!", "http://apice.unibo.it/xwiki...")
15     Web2 = FakeSearch("web2", "Distributed Systems are the best!", "http://apice.unibo.it/xwiki...")
16     Image1 = FakeSearch("image1", "But what's really a distributed system?", "https://www.unibo.it/uniboweb...")
17     Image2 = FakeSearch("image2", "But what's really a distributed system?", "https://www.unibo.it/uniboweb...")
18     Video1 = FakeSearch("video1", "Cool video about distributed systems", "https://www.youtube.com/watch?v=...")
19     Video2 = FakeSearch("video2", "Cool video about distributed systems", "https://www.youtube.com/watch?v=...")
20 )
```

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - **Serial request**
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

# googleSerial.go

## DIY

- Let's implement a function that will retrieve results from the different repositories (Web, Images, Videos).
- We should fill an array on type `Result` with the results of each of these searches.

### Serial search function's signature

```
1 func SearchSerial(query string) ([]Result, error) {  
2     // TODO should return an array of results made of every possible result from Web, Image, Video for a specific query  
3     return nil, nil  
4 }
```

# googleSerial.go

## SOLUTION

- By being serial, this function has serious scalability problems: each search on each repository will only begin when the previous one has finished
- What if a search will take a long time to complete?
- The system will be stuck and the user unhappy
- NB: keep an eye on the amount of elapsed time. We're going to reduce that a lot!

```
1 func SearchSerial(query string) ([]Result, error) {  
2     results := []Result{  
3         Web(query),  
4         Image(query),  
5         Video(query),  
6     }  
7     return results, nil  
8 }
```

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - **Parallel request**
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system



# googleParallel.go

## DIY

- Let's iterate over our past serial solution and let's make it parallel by using Channels and Goroutines!
- The search will be still run on a single machine, but by using Goroutines we'll be able to search in every repository all at once
- We're still bound to each search function to return to our main control flow. Only after they all returned we can finally move on

```
1 func SearchParallel(query string) ([]Result, error) {  
2     // should return the same result as the SearchSerial function, but the search should be parallelized  
3     // TODO create a channel of type Result and use Goroutines to parallelize the work you did in SearchSerial  
4     // HINT send each search result on the channel and receive those results inside the array of results  
5     return []Result{}, nil  
6 }
```

# googleParallel.go

## SOLUTION

- Look at the elapsed time! It has been reduced by 1/3
- Also note how little more code we had to add to the previous solution in order to get this much benefit
- This solution still has problems though
- The elapsed time will now be capped to the slowest search that has been run amongst all
- Let's try to iterate once again and make it better

```
1 func SearchParallel(query string) ([]Result, error) {  
2     c := make(chan Result)  
3     go func() { c <- Web(query) }()  
4     go func() { c <- Image(query) }()  
5     go func() { c <- Video(query) }()  
6  
7     return []Result{<-c, <-c, <-c}, nil  
8 }
```

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - **Timeout-proof request**
  - Replicated/distributed request
  - Embed a Webserver in the system

# googleTimeout.go

## DIY

- We want to be able to get as many results as possible inside a predefined window of time, so that the elapsed time will always be shorter than X ms (timeout), no matter the query
- Consider the timer as an Agent that sends a message on the timer channel whenever it expires. Therefore, you can use it just like you use channels you defined yourself

```
1 func SearchTimeout(query string, timeout time.Duration) ([]Result, error) {  
2     // should return the same result in the same way as the SearchParallel function, but only if a timeout doesn't occur  
3     // TODO use the timer and the select statement inside the result array to either receive the search results from  
4     // the channel you created before, or return with what you have when the timer expires  
5     // HINT you can think of the timeout as a message on a whatever channel, also return a proper error string  
6     timer := time.After(timeout)  
7  
8     return nil, nil  
9 }
```

# googleTimeout.go

## SOLUTION

- Look at the next slide for further explanation

```
1 func SearchTimeout(query string, timeout time.Duration) ([]Result, error) {
2     timer := time.After(timeout)
3     c := make(chan Result, 3)
4
5     go func() { c <- Web(query) }()
6     go func() { c <- Image(query) }()
7     go func() { c <- Video(query) }()
8
9     var results []Result
10    for i := 0; i < 3; i++ {
11        select {
12            case result := <-c:
13                results = append(results, result)
14            case <-timer:
15                return results, errors.New("timed out")
16        }
17    }
18    return results, nil
19 }
```

# googleTimeout.go

## SOLUTION

- Now our search function will either return all the results if they took less than the timeout amount of time, or just the results it got before it expired
- Try and play with the timeout value, to see what happens if it becomes too little or very big
- The main problem left is that the search function is run on a single machine, on a single location for all the data.
- What if we could have multiple machines, and i.e. the closest ones to the source of the request could answer faster than the others?

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - **Replicated/distributed request**
  - Embed a Webserver in the system

# googleReplicated.go

## DIY

- Let's solve the replication problem!
- We will first need to modify `SearchTimeout` in order to use the `replicated*FakeSearchFunction` that we have in our `searchUtils.go` file

```
1 func SearchReplicated(query string, timeout time.Duration) ([]Result, error) {  
2     // TODO same as you did in timeout, but this time use the replicated*FakeSearchFunctions  
3     // TODO also go complete the FirstFetchAmong function in searchUtils  
4     return nil, nil  
5 }
```



# googleReplicated.go

## SOLUTION

- Here's the solution to the previous task
- As you can clearly see, not much has changed from SearchTimeout, apart from the functions we call to retrieve the results from the various repositories

```
1 func SearchReplicated(query string, timeout time.Duration) ([]Result, error) {
2     timer := time.After(timeout)
3     c := make(chan Result, 3)
4     go func() { c <- replicatedWebFakeSearch(query) }()
5     go func() { c <- replicatedImageFakeSearch(query) }()
6     go func() { c <- replicatedVideoFakeSearch(query) }()
7
8     var results []Result
9     for i := 0; i < 3; i++ {
10         select {
11             case result := <-c:
12                 results = append(results, result)
13             case <-timer:
14                 return results, errors.New("timed out")
15         }
16     }
17     return results, nil
18 }
```

# fetchFirstAmong.go

## DIY

- Now let's implement a new function that will be used to distribute a search among all the machines we have
- In this example the different machines will just be simulated, but it should give you an idea on how it should work in practice

```
1 func FirstFetchAmong(replicas ...SearchFunc) SearchFunc {  
2     return func(query string) Result {  
3         // TODO create a buffered channel of type Result and as long as the number of replicas that were passed in input  
4         // TODO create an inner function that given a index runs the replicated function at that index of replicas  
5         // TODO launch this function in a new goroutines for every index of replicas  
6         // TODO return the very first value received on the channel  
7         return nil  
8     }  
9 }
```

# fetchFirstAmong.go

## SOLUTION

- Now whenever we're triggering a new query the system will handle it to different instances, and return the various results from whichever machine solved it faster
- For example, for a certain search we could get the Web(query) result from Machine #1, and the Images(query) result from Machine #2.

```
1 func FirstFetchAmong(replicas ...SearchFunc) SearchFunc {  
2     return func(query string) Result {  
3         c := make(chan Result, len(replicas))  
4         searchReplica := func(i int) {  
5             c <- replicas[i](query)  
6         }  
7         for i := range replicas {  
8             go searchReplica(i)  
9         }  
10        return <-c  
11    }  
12 }
```

# Outline

- 1 The GO programming language
  - Introduction to GO
  - PROs and CONs of GO
  - Install GO on your machine
  - Basic syntax and constructs
  - Useful idiomatic language features
- 2 GO in practice with a Concurrency exercise
  - Ex: Ping Pong
- 3 Building a distributed system step by step
  - Introduction to the system
  - Serial request
  - Parallel request
  - Timeout-proof request
  - Replicated/distributed request
  - Embed a Webserver in the system

## ws.go

```

1 func main() {
2     http.HandleFunc("/fakeGoogleSearch", handleSearch) // init a new API
3     fmt.Println("serving on http://localhost:8080/fakeGoogleSearch")
4     log.Fatal(http.ListenAndServe("localhost:8080", nil))
5 }
6
7 func handleSearch(w http.ResponseWriter, req *http.Request) {
8     log.Println("serving", req.URL)
9     start := time.Now()
10    results, err := google.SearchReplicated("golang", 104*time.Millisecond) // run a new search
11    elapsed := time.Since(start)
12    if err != nil {
13        http.Error(w, err.Error(), http.StatusInternalServerError)
14        return
15    }
16    resp := google.Response{Results: results, Elapsed: elapsed}
17    responseTemplate.Execute(w, resp) // build and print to screen the results
18    fmt.Println(resp, err)
19 }
20
21 var responseTemplate = template.Must(template.New("results").Parse(
22     '<html><head/><body>
23     <ol>{{range .Results}}
24         <li>{{.Title}} - <a href="{{.URL}}">{{.URL}}</a></li>
25     {{end}}</ol>
26     <p>{{len .Results}} results in {{.Elapsed}}</p>
27 </body></html>'))

```

# Final step

## The webserver

What we saw in the last slide is a simple implementation of a WebServer in GO that uses what we built so far, and prints it on screen, on a webpage that (with a little fantasy on our part) kind of reminds us of Google!





```
1. Query: "golang" | Source: web2 -> Result: Distributed Systems are the best! - http://apice.unibo.it/xwiki/bin/view/Courses/Sd17/  
2. Query: "golang" | Source: image2 -> Result: But what's really a distributed system? - https://www.unibo.it/uniboweb/utills/UserIn  
3. Query: "golang" | Source: video1 -> Result: Cool video about distributed systems - https://www.youtube.com/watch?v=dQw4w9
```

3 results in 101.20257ms

We can now say that what we built is an Hybrid Architecture[3] (Horizontal and Vertical):

- Horizontal: there are many replicated machines computing a certain task (a specific search in our case)
- Vertical: there's a machine that acts like a middleware between the client and the  $n$  machines that should compute a certain task.

# Bibliography

-  E. 800. “Definitions of terms related to quality of service”. In: *ITU-T Recommendation* (1994).
-  Sameer Ajmani. *Program your next server in GO*. 2016. URL: <https://www.youtube.com/watch?v=5bY060-qY0I>.
-  Jmaarten Van Steen Andrew S. Tanenbaum. *Distributed Systems, Principle and Paradigms*. PHI, 2009.
-  Rob Pike. *Less is exponentially more*. 2012. URL: <https://commandcenter.blogspot.it/2012/06/less-is-exponentially-more.html>.